

Gen AI Model Research													
Sl No	Provider	Model Name	Release Date	Knowledge Cutoff	Input Cost (\$/1M tokens)	Output Cost (\$/1M tokens)	Context Window	Model Type	Open / Closed	Multimodal?	Best Use Case	Strengths	Limitations
1	OpenAI	GPT-5.5	April 2026	Dec 2026	\$5.00	\$30.00	1.05M tokens	Foundation / Multimodal	Closed	Yes (text, image, audio)	autonomous, agentic tasks, advanced coding, and complex, multi-step knowledge work	stronger coding performance and more reliable tool use. It stays on task for significantly longer without stopping early, which matters most for the complex, long-running work our users delegate to Cursor.	Strict Usage Limits, High Costs for Long Context,Agentic Orchestration Necessity, Performance/Cost Trade-off: High-effort reasoning (ahigh) to solve complex bugs can rapidly exhaust usage quotas.
2	Google Deepmind	Gemini 3.1	Feb 2026	Jan 2026	\$2.00	\$12.00	2M tokens	Multimodal/ Agentic	Closed	Yes (text, image)	Complex problem-solving across text, audio, images, video, PDFs, and entire code repositories; agentic workflows, finance, and spreadsheet applications	Largest readily-available context (1M); native video+audio; Google ecosystem synergy; competitive pricing	Latency increases with very long contexts; multimodal adds compute overhead; less strong on pure text benchmarks
3	Anthropic	Claude Opus 4.7	Apr 2026	Jan 2026	\$5.00	\$25.00	1M tokens	Frontier reasoning/Agentic model	Closed	Yes (text, image)	Long-running, asynchronous agents; large codebases, multi-stage debugging, end-to-end project orchestration, and knowledge work like drafting documents and analyzing data	Stronger agentic coding, self-verifying outputs, more literal instruction-following, 3.3x higher-resolution vision, and improved file-system memory across multi-session agent work	A new tokenizer can produce up to 35% more tokens for the same input text, meaning effective costs per request can be higher than the rate card suggests; Finout slower than average and very verbose compared to similarly priced models
4	Meta	Llama 4 Scout	April 2026	Mar 2026	\$0.15	\$0.30	10M tokens	Open-Weight	Open	Yes (text, image)	Cost-sensitive large-scale deployments, research, fine-tuning for custom domains, self-hosted AI	World's largest context window (10M tokens!); open weights for self-hosting & fine-tuning; extremely cheap via API	High Hardware Requirements, open weights raise misuse concerns,Licensing Constraints; Image generation constrains
5	xAI	Grok 4.3	Apr 2026	Nov 2026	\$1.25	\$2.50	1M tokens	flagship Reasoning/Agentic workflow	Closed	Yes (text, image, video)	Real-time X/Twitter data tasks, up-to-date news Q&A, complex reasoning chains,Multimodal Video & Image Analysis	Real-time X (Twitter) data access; strong coding	Tied to X ecosystem; less enterprise adoption; safety guardrails less mature than OpenAI/Anthropic
6	DeepSeek	DeepSeek V4	Apr 2026	Apr 2026	\$0.44	\$0.87	1.045M tokens	Reasoning / MoE (mixture of experts)	Open	No (text only)	Math & scientific reasoning, STEM education, cost-efficient chain-of-thought tasks, open-source research	90%+ on math benchmarks; open-weight (671B MoE, 37B active); cheaper than GPT-4; transparent reasoning traces	Text-only (no images/video); smaller context (128K); some benchmark gaps vs frontier closed models; Chinese company privacy considerations
7	Alibaba	Qwen 3.6 Plus (Flagship)	Apr 2026	Early 2026	\$0.33	\$1.95	1M tokens	MoE	Closed	Yes (text, image, audio)	Agentic coding, front-end development, 3D scenes, games, repository-level problem solving, and complex reasoning tasks	very competitive pricing,Very high output speed	High Fabrication and Hallucination Rates,Limitations in Specific Languages and Syntax,"Loopy" Agentic Behavior,Long-Context Inconsistency
8	Mistral AI(France)	Mistral Large 3	Dec 2026	NA	\$0.50	\$1.50	262K tokens	MoE	Open	Yes	Production-grade assistants, retrieval-augmented generation (RAG) systems, scientific workflows, complex enterprise workflows, and long-document analysis	61.6% SWE-bench Verified — beats Gemini 2.5 Pro & GPT-4.1 at ¼ the price; open-weight; GDPR-friendly for EU orgs	Not a dedicated reasoning model — dedicated reasoning models can outperform it in strict reasoning tasks; lags behind vision-first models on multimodal tasks; complex to deploy efficiently at scale due to its large size
9	Moonshot AI (China)	Kimi K2.6	Apr 2026	Apr-25	\$0.60	\$2.50	262K tokens	Foundation / MoE	Open	Yes (text, image, audio)	Long-horizon coding, coding-driven design, proactive autonomous execution,swarm-based task orchestration — delivering documents, websites, and spreadsheets in a single autonomous run	12-hour continuous, autonomous task execution,specialized coding capabilities	Substantial Hardware Requirements,Context Window Limitations,Inconsistent System-Level Understanding,Lower Accuracy on Simple Tasks
10	Google DeepMind	Gemma 3	March 2026	Aug 2026	\$0.00	\$0.13	128K tokens	Code / Agentic	Open	Yes (text, image)	Runs 100% locally — no API keys, no data leaks, safe for proprietary code,Native function calling — no prompt hacking needed for tool use,Zero per-token cost — run unlimited tests locally for free	Runs on consumer-grade GPUs like an NVIDIA RTX 3090 with quantization; supports over 140 languages; built on the same research as Gemini	Zero per-token cost — run unlimited tests locally for free,Text output only, Creative tasks (UI design, test report generation) are weak,Self-hosting required
Notes:													
1 Need the best result, cost doesn't matter → Claude Opus 4.7 or GPT-5.5													
2 Need multimodal (video/audio/image) → Gemini 3.1 Pro													
3 Building at scale, budget is critical → DeepSeek V4 or Qwen 3.6													
4 Data must stay private / on-premise → Gemma 4 or Mistral Large 3													
5 Long autonomous coding runs → Kimi K2.6													
6 Massive codebase, millions of tokens → Llama 4 Scout													
7 Real-time data / X/Twitter context → Grok 4.3													
8 EU compliance + multilingual → Mistral Large 3													